

SATVIC MOVEMENT • FOUNDER'S OFFICE • AI

Reaching the India that Hindi can't

Dubbing our long-form library into regional languages — in Subah's own voice, using AI.

2

languages — Telugu & Malayalam

0

voice-leaks across 460 segments

₹0

spent of the ₹2,500 budget

THE OPPORTUNITY

Most of India isn't watching in Hindi

96M

Telugu

largest single unreached audience

38M

Malayalam

Kerala — deeply health-conscious

59M

Kannada

next in line (Bengaluru market)

≈**190 million people** speak these three languages — and the channel is invisible to them today, purely for lack of an audio track.

WHY NOW

YouTube dubs for free — just not into these

What YouTube gives away free

- ✓ Auto-dubbing from Hindi into 16 languages
- ✓ English, Bengali, Punjabi, Spanish, Arabic...
- ✓ Multi-audio tracks on one video

...but none of this reaches Telugu / Malayalam / Kannada.

The gap we fill

- Hindi → Telugu / Malayalam / Kannada: not supported
- Kannada isn't an auto-dub target from any language
- Where auto-dub exists, it's a generic robot voice — not Subah's

Custom voice-cloned tracks are the only route — and almost no Hindi-first channel has taken it.

THE RESULT

What we shipped

2

dubbed videos

Telugu + Malayalam, 16:12 each

100%

Subah's voice

zero-shot clone from ~7s of her audio

0 / 460

voice-leaks

verified by content-based scan

₹0

tool cost

fully local, open-source stack

-14 LUFS

broadcast loudness

music bed preserved under the dub

100%

reviewable

every line back-translated for QC

HOW IT WORKS

The pipeline — open-source, end to end

1 INGEST

yt-dlp • Whisper • Demucs

- Pull the video
- Timestamped Hindi transcript
- Split voice / music bed

2 TRANSLATE

LLM + glossary + human

- Glossary-locked terms
- Everyday colloquial register
- Back-translated • native-reviewed

3 VOICE

IndicF5 (AI4Bharat)

- Clone Subah from ~7s
- Synthesize each segment
- Leak-scan + trim → 0 leaks

4 ASSEMBLE

ffmpeg

- Fit each line to timing
- Remix original music bed
- Mux → YouTube audio track

Translation lives outside the dubbing tool on purpose — for health content, every word is frozen and reviewed before any audio exists.

TOOL SELECTION

Why IndicF5 — and the two pivots to get there

Sarvam AI Dub

Purpose-built for Indic dubbing. Won on paper — but the voice-clone product is waitlist-gated.

PIVOTED AWAY

ElevenLabs

Best clone fidelity; Telugu + Malayalam supported. ₹1,950 — kept as a priced fallback, never needed.

HELD IN RESERVE

IndicF5 (AI4Bharat)

IIT-Madras open-source. Voice clone + native Te/MI, runs locally, no watermark, no subscription.

CHOSEN • ₹0

The ₹2,500 budget stopped being a spend and became an escalation ladder we priced but never had to climb.

ACCURACY

Fluent isn't enough — it has to be right

- 1 Glossary lock** Brand + ayurvedic terms, English loanwords, numbers as words — decided before translation.
- 2 Back-translation** Every one of 460 lines carries a literal English back-translation for audit.
- 3 Health-claim flags** 74 Telugu · 63 Malayalam lines pre-flagged: dosages, durations, do/don't.
- 4 Native watch-down** A native speaker reviews the script, then watches the full dub. This is the sign-off.
- 5 Technical QC** Duration, sync, -14 LUFS loudness, music bed, zero voice-leaks.

The native review is the gate — not the machine. A smooth-sounding wrong dub about someone's health is worse than no dub.


WHAT DIDN'T WORK

The bug a native ear caught — twice

IndicF5 occasionally echoes its voice-reference into a segment's start. A first "neutral" reference just changed which words leaked — and because they were English, an automated scan reported a false "0." The listener caught it.

Root cause: a badly-cut reference clip. A clean, complete-sentence reference fixed it at the source.

Reference-echo leak rate (% of segments)



The lesson A detector's green light means nothing until it's stress-tested for false positives AND false negatives. The human ear is the ground truth — which is exactly why native review is built into the process, not bolted on.

COST

₹0 spent. The budget became insurance.

₹0

of the ₹2,500 tool budget

Every tool free & local: yt-dlp · Whisper · Demucs · IndicF5 · ffmpeg

The priced escalation ladder — never needed

Tier 1 — IndicF5 (local)

SHIPPED

₹0

Tier 2 — ElevenLabs Starter

if reviewers want more polish

≈₹450

Tier 3 — ElevenLabs Creator

max fidelity, comfortable headroom

≈₹1,950

From one video to the whole library

1

Pipeline is a queue

Everything above is scripted. Machines do the repeatable work; humans only touch script-review and the native watch-down.

2

Order by analytics

Dub top watch-time videos first; pick languages from Geography data — not intuition.

3

The glossary compounds

Every reviewed video grows the locked glossary, so per-video review time keeps shrinking. It's the moat.

4

Keep auto-dub free

Let YouTube cover the 16 languages it does support from Hindi; spend effort only where it's blind — Te / MI / Kn / Ta.

5

Cost is QC, not compute

One GPU collapses synthesis to minutes. The real cost is native review (~₹2-3k/video) — where a health channel should spend.

6

Formalize the voice

One consented, verified clone per language in a Satvic-controlled account — consistency + clean rights.

A repeatable system, not a one-off.

Two clean dubs today, in Subah's voice, for ₹0 — and a pipeline ready for the full library.

NEXT STEPS

- Native-speaker sign-off on both dubs (review kit ready)
- Pick the next 10–20 videos by YouTube Geography analytics
- Record Subah's voice-consent; stand up one clone per language